


Introduction to Neural Networks

David, Michael, Teddy, Tommy





What is a Neural Network

- Machine Learning
 - Regression
 - Classification
- Applications
 - Image Classification
 - Speech Recognition
 - Large Language Models




Neural Network Architecture







$$\hat{y} = wx$$




$$\hat{y} = wx$$




$$\hat{y} = wx$$







$$\hat{y} = wx + b$$




$$\hat{y} = wx + b$$




$$\hat{y} = wx$$




$$\hat{y} = wx + w_b$$







$$\hat{y} = wx + w_b$$




$$\hat{y} = w_1 x_1 + w_2 x_2 + w_b$$






$$\hat{y} = w_1 x_1 + w_2 x_2 + w_b$$




Hidden Layer



$$N_1 = w_{11}x_1 + w_{21}x_2 + b_{N1}$$

$$N_2 = w_{12}x_1 + w_{22}x_2 + b_{N2}$$

$$\hat{y} = w_{1y}N_1 + w_{2y}N_2 + b_y$$






$$N_1 = \sigma(w_{11}x_1 + w_{21}x_2 + b_{N1})$$


$$N_2 = \sigma(w_{12}x_1 + w_{22}x_2 + b_{N2})$$

$$\hat{y} = \sigma(w_{1y}N_1 + w_{2y}N_2 + b_y)$$

Activation Functions



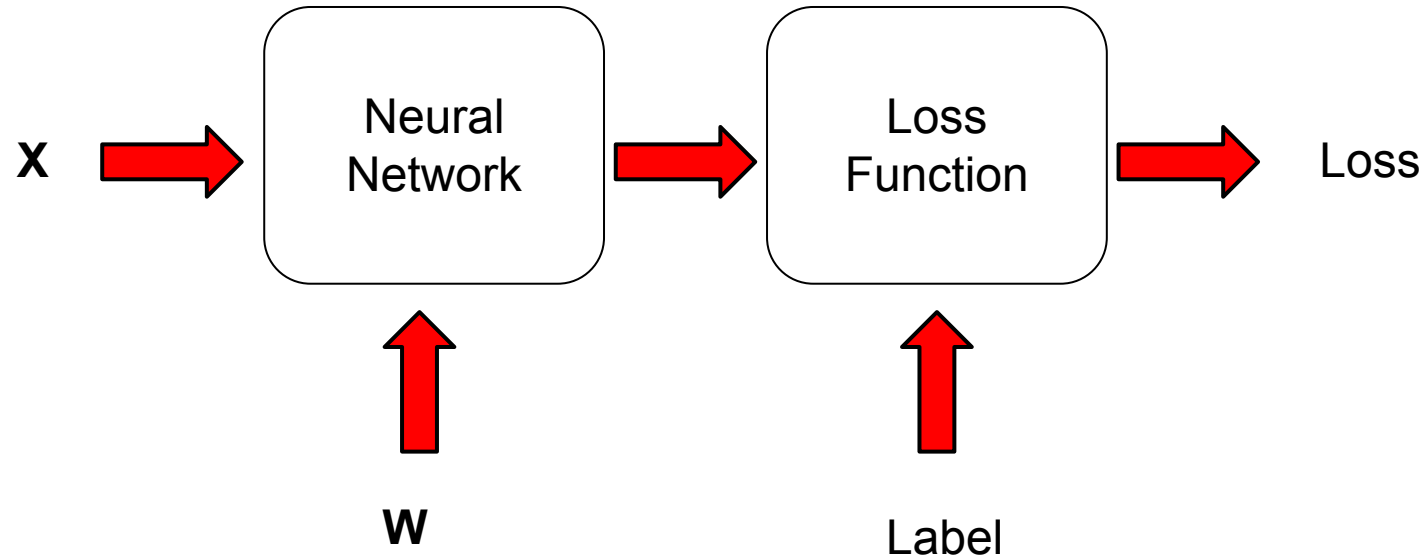
See <https://keras.io/api/layers/activations/> for more common activation functions



Training Neural Networks


Loss Functions

- Regression
 - Mean Squared Error
 - Mean Absolute Error
- Classification
 - Cross entropy
- See <https://keras.io/api/losses/> for more examples




$f(\mathbf{W})$: Loss as a function of weights for a given **X**


Gradient Descent




Gradient Descent




Gradient Descent



Gradient Descent



Gradient Descent



Optimizers

- Stochastic Gradient Descent (SGD)
- Adam
- RMSProp
- See <https://keras.io/api/optimizers/> for other common optimizers